

Week 8 - 6103 Handout

Subhayan Mukerjee

10/3/2021

Welcome to R!

This document is an RMarkdown document. It allows you to write text as well as code (and their outputs) in a single convenient file, and is often used for tutorial purposes. Most of what you will be doing in R, however, will be done in a R script file. An R Script file is essentially a text file, with the .R extension. You can create an R Script file, by opening RStudio and then clicking on File -> New File -> R Script. Alternatively, on Windows you can simply open RStudio and press Ctrl + Shift + N.

When you open RStudio you will see there are four quadrants as follows (Figure 1)

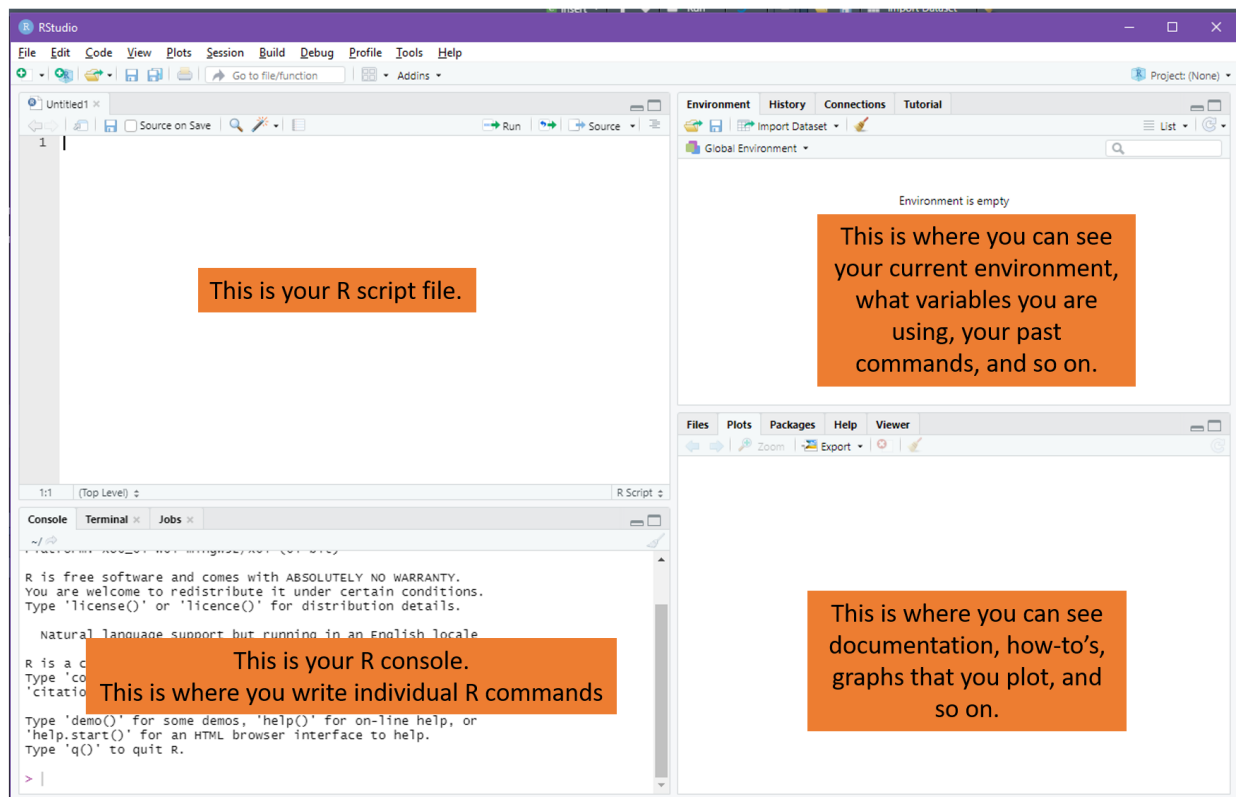


Figure 1: RStudio

Let's type our first R command. Take your cursor to the console (bottom left) and type the following (what you see in the shaded section). This command simply tells R to display the text "Hello World!". You should see the output that follows.

```
print("Hello World!")
```

```
## [1] "Hello World!"
```

In this document, like in all RMarkdown documents, the R commands appear on a shaded background. And the output from the R command appears right after the ‘##’. The [1] in the output refers to the first line of output.

You can do basic arithmetic in the console as well. For example

```
1+1
```

```
## [1] 2
```

```
5-2
```

```
## [1] 3
```

```
22*45/56
```

```
## [1] 17.67857
```

```
sqrt(98)
```

```
## [1] 9.899495
```

```
5^3
```

```
## [1] 125
```

However, as you will notice, you can only write single line commands in the console. As most analysis that you will do, will require multiple commands, it is wiser to use the R Script (top left)

Type out all the commands that you just typed in the console, in the R script, select the entire text and click on “Run” (Figure 2). The output will appear in the console.

You can also select each line individually and click on “Run” to run just that line.

From now on, when you encounter an R command in this file, type it out in the R script and run it from there, instead of from the console.

Variables

Now, as you know we can’t keep dealing with numbers. We will need to store them in variables. These variables are not conceptually the same as social science variables we have discussed so far, but they can be used as a proxy. You can use R variables to store actual variables that you have collected for your research.

Creating variables in R is easy. For example, the next command defines a variable called ‘a’ and assigns to it, the value of 5. You can use both, the “=” as well as the “<=” symbols to assign values to a variable.

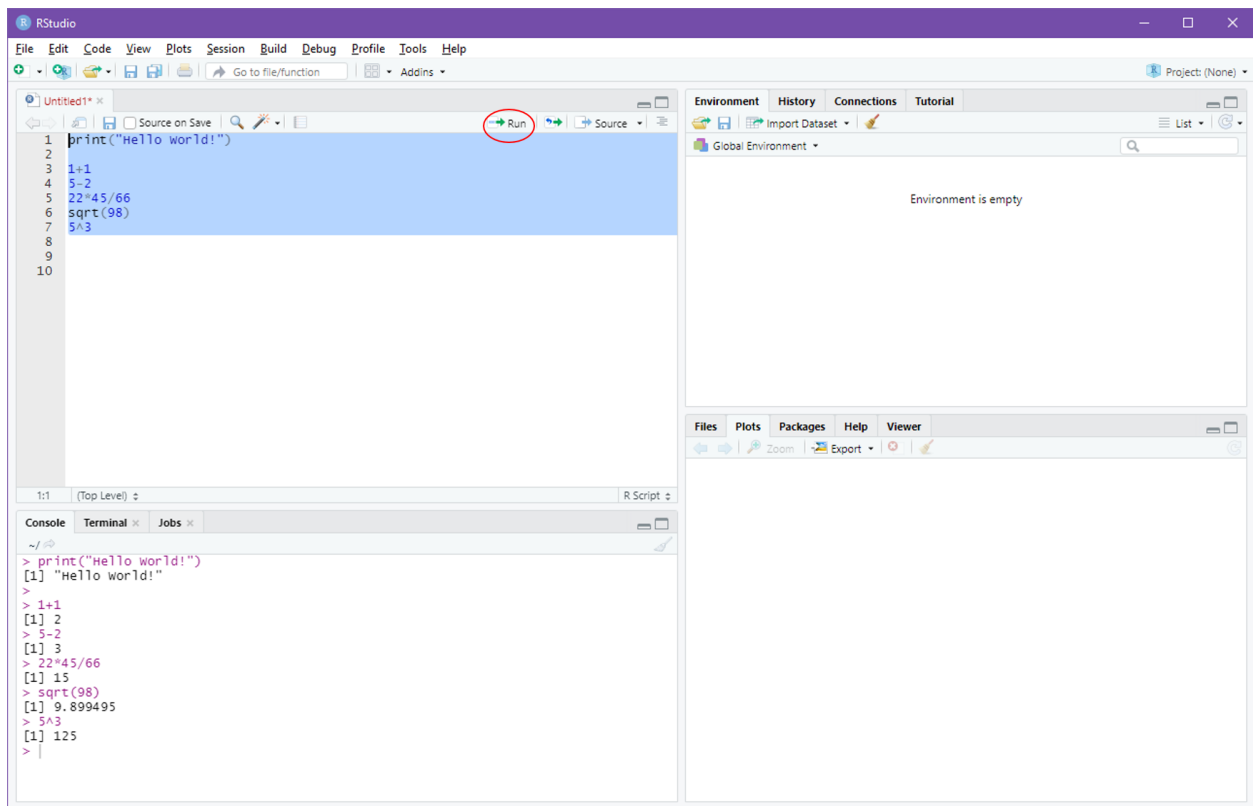


Figure 2: R script in R Studio

```
a <- 5
```

Once a variable has been assigned a value, it will appear in your environment panel of RStudio (top right) (Figure 3).




Figure 3: R script in R Studio

you can see what the value is by doing:

```
print(a)
```

```
## [1] 5
```

Variables are just like numbers. You can do arithmetic on them.

```
a <- 6
b <- 4
c <- 5

d <- (a + b) / c

print(d)
```

```
## [1] 2
```

Vectors

Vectors lets you combine multiple values into a single variable. Say, for instance, we have the heights of 5 individuals (in cm.). We could have a single variable for each of them (`h1`, `h2`, `h3`, `h4`, `h5`) but that can be very confusing to keep track of. Instead you can define a vector called heights as follows:

```
heights <- c(140, 176, 154, 143, 166, 180, 148, 178)

print(heights)
```

```
## [1] 140 176 154 143 166 180 148 178
```

The `c` in the above command stands for “combine”. It instructs R to combine the five values into one vector. Remember that vectors are also “variables”. Anything that can take a value is a variable. Vectors are just variables that are made up other variables.

We can access individual elements of a vector using square brackets. The following two commands gives you the first and third height respectively.

```
print(heights[1])
```

```
## [1] 140
```

```
print(heights[3])
```

```
## [1] 154
```

The number within the square brackets is called the index. These individual elements are similar to the variables (`a`, `b`, `c`) we used earlier. Indeed, you can combine existing variables into a vector as well like this:

```
random_vector <- c(a, b, c, d)

print(random_vector)
```

```
## [1] 6 4 5 2
```

Character variables

Similar to numeric variables you can also define textual variables. These are called **character** variables. You can have vectors of these variables as well.

```
text_var <- "abcdef"
text_var2 <- "xyz"

print(text_var)
```

```
## [1] "abcdef"
```

```
print(text_var2)
```

```
## [1] "xyz"
```

```
persons <- c("Michael", "Alice", "Sanjay", "Claudia", "Bob", "Jack", "Mary", "Jill")  
print(persons[4])
```

```
## [1] "Claudia"
```

```
print(persons[2])
```

```
## [1] "Alice"
```

What happens when you try printing the 10th person's name (that clearly doesn't exist?)

```
print(persons[10])
```

```
## [1] NA
```

Logical variables

Logical variables can either be TRUE or FALSE. They are used to know whether a certain variable is equal to (== NOTE the double equals), less than (<), or greater (>) than a certain value. These can also be combined as “less than or equal to” (<=) or “greater than or equal to” (>=) Let's look at an example:

```
print(a)
```

```
## [1] 6
```

```
print(b)
```

```
## [1] 4
```

```
print(c)
```

```
## [1] 5
```

```
print(a == 6) # is a equal to 6?
```

```
## [1] TRUE
```

```
print(b > 3) # is b greater than 3?
```

```
## [1] TRUE
```

```
print(c <= 5) # is c less than or equal to 5?
```

```
## [1] TRUE
```

What you see on the right of each command, after the `#` is called a comment. They are used to describe what that command does, and is ignored by R. They are largely written to make the script understandable to someone reading the file.

Another useful operator is the `!=` which means “not equal to”. It can be used as follows

```
print((a+b) != 10) # is a+b not equal to 10?
```

```
## [1] FALSE
```

Logical variables are useful in many cases. Let’s look at the `heights` vector again.

```
print(heights)
```

```
## [1] 140 176 154 143 166 180 148 178
```

We want to know which of the heights are greater than 160?

```
print(heights > 160)
```

```
## [1] FALSE TRUE FALSE FALSE TRUE TRUE FALSE TRUE
```

This returns a vector of logicals. If you want to know what those heights values are (corresponding to `TRUE`), you can simply use the output of the above command as the index for the vector. The way R works, it only shows those values which correspond to `TRUE`.

```
print(heights[heights > 160])
```

```
## [1] 176 166 180 178
```

This command essentially prints all the heights that are above 160. What does the next command do?

```
print(heights[heights == 170])
```

```
## numeric(0)
```

Dataframe

A dataframe is a table, comprising of rows and columns of data. They are the most widely used data structures in R. Let’s put the two vectors we have into a dataframe. To do that, we will first use the `cbind` or the column bind command to put the two vectors “next to each other” as two columns, and then use the `data.frame` command to create a dataframe.

Finally, we will assign the dataframe to a variable called `df` (the name can be anything, obviously).

This gives us a dataframe with two columns and 8 rows. Each row corresponds to one person.

```
df <- data.frame(cbind(persons, heights))
print(df)
```

```
##  persons heights
## 1 Michael     140
## 2   Alice     176
## 3  Sanjay     154
## 4 Claudia     143
## 5    Bob     166
## 6   Jack     180
## 7   Mary     148
## 8   Jill     178
```

We can access the individual columns using the `$` command:

```
print(df$persons)
```

```
## [1] "Michael" "Alice"  "Sanjay"  "Claudia" "Bob"    "Jack"    "Mary"
## [8] "Jill"
```

```
print(df$heights)
```

```
## [1] "140" "176" "154" "143" "166" "180" "148" "178"
```

What each of these two commands gives you are the original vectors that you used to make the dataframe in the first place.

However, if you notice closely, you will see that `df$heights` shows you the height values but within double quotes. This is because when you used `cbind` it converted both vectors into a common format. In other words, the heights were converted from numeric to character. If you're not sure what I mean by this use the `class` command. The `class` command tells you the type of the data stored in that variable. So, run the following two commands:

```
class(heights)
```

```
## [1] "numeric"
```

```
class(df$heights)
```

```
## [1] "character"
```

`heights` is numeric but `df$heights` is a character.

So we will need to reconvert them to numeric. So use the `as.numeric` command as shown below, and re-assign the re-converted heights to the same column.

```
df$heights <- as.numeric(df$heights)
```


Dataframes are very powerful as they allow you to do a variety of analysis. Crucial is knowing how to index dataframes, for retrieving the information you need.

Just like vectors are indexed using square brackets, so are dataframes. But unlike vectors they need two numbers, one for the row, and one for the column.

For example

```
print(df[2,1])
```

```
## [1] "Alice"
```

```
print(df[6,2])
```

```
## [1] 180
```

Again, what happens when you index incorrectly?

```
print(df[3,4])
```

```
## NULL
```

Now these indices can be logicals. This lets us ask questions like, “who are the persons with heights greater than 160?”

Let’s answer this step-by-step.

First, we have

```
print(df$heights > 160)
```

```
## [1] FALSE  TRUE FALSE FALSE  TRUE  TRUE FALSE  TRUE
```

Next, we use this as the first index for our dataframe (we want the rows where heights > 160). For the second index, we keep it blank, so that we can see all (both) columns.

```
print(df[df$heights > 160,])
```

```
##  persons heights
## 2   Alice      176
## 5    Bob      166
## 6   Jack      180
## 8   Jill      178
```

Remember that the previous output is a dataframe itself. So, if we just want the heights, we can do

```
print(df[df$heights > 160,]$heights)
```

```
## [1] 176 166 180 178
```

You can also sort the dataframe in ascending or descending order of height using the `order` command:

```
print(df[order(heights),]) # increasing (ascending) order
```

```
## persons heights
## 1 Michael      140
## 4 Claudia      143
## 7 Mary         148
## 3 Sanjay       154
## 5 Bob          166
## 2 Alice        176
## 8 Jill         178
## 6 Jack         180
```

```
print(df[order(-heights),]) # decreasing (descending) order
```

```
## persons heights
## 6 Jack         180
## 8 Jill         178
## 2 Alice        176
## 5 Bob          166
## 3 Sanjay       154
## 7 Mary         148
## 4 Claudia      143
## 1 Michael      140
```

Descriptive Statistics

Now let's look at some statistics of the heights

```
print(mean(df$heights)) # this gives the mean height
```

```
## [1] 160.625
```

```
print(median(df$heights)) # this gives the median height
```

```
## [1] 160
```

```
print(sd(df$heights)) # this gives the standard deviation of the heights
```

```
## [1] 16.39632
```

Other commands you can use are `min` (for minimum) or `max` for maximum. You can see a statistical summary of a variable using the `summary` command:

```
summary(df$heights)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  140.0   146.8   160.0   160.6   176.5   180.0
```

Let's add another column to the dataframe, let's call it gender.

```
df$gender <- c("M", "F", "M", "F", "M", "M", "F", "F")
print(df)
```

```
##   persons heights gender
## 1 Michael     140      M
## 2   Alice     176      F
## 3  Sanjay     154      M
## 4 Claudia     143      F
## 5    Bob      166      M
## 6   Jack      180      M
## 7   Mary      148      F
## 8   Jill      178      F
```

Let's now find the average heights of Males

Step 1: select only those rows where gender == "M" (Remember the double equals "==")

```
print(df[df$gender == "M",])
```

```
##   persons heights gender
## 1 Michael     140      M
## 3  Sanjay     154      M
## 5    Bob      166      M
## 6   Jack      180      M
```

Step 2: get their heights

```
print(df[df$gender == "M",]$heights)
```

```
## [1] 140 154 166 180
```

Step 3: Calculate their mean

```
print(mean(df[df$gender == "M",]$heights))
```

```
## [1] 160
```

Repeat for females

```
print(mean(df[df$gender == "F",]$heights))
```

```
## [1] 161.25
```

You can inspect a dataframe anytime by using the `str` or `structure` command.

```
str(df)
```

```
## 'data.frame':   8 obs. of  3 variables:
## $ persons: chr  "Michael" "Alice" "Sanjay" "Claudia" ...
## $ heights: num  140 176 154 143 166 180 148 178
## $ gender : chr  "M" "F" "M" "F" ...
```

Now let's save the data that we generated in a CSV (comma separated values) file

```
write.csv(df, "heights.csv")
```

Where has the file been saved? In your current working directory. Type the command `getwd()` to locate it.

Let's work on an actual dataset!

Download the `gapminder.csv` file from [luminus](https://luminus.net). Let's first read in the file into R.

```
gapminder_df <- read.csv("data/gapminder.csv") # replace this path with where you downloaded the file
```

We want to look at the data, but this dataset has thousands of rows, so we don't want to see all of them. If you do `print(gapminder_df)` R will only show you the first 1000 rows in the console. You can try that, but it's not really very helpful. Let's look at the structure instead:

```
str(gapminder_df)
```

```
## 'data.frame':    1704 obs. of  6 variables:
## $ country   : chr  "Afghanistan" "Afghanistan" "Afghanistan" "Afghanistan" ...
## $ continent: chr  "Asia" "Asia" "Asia" "Asia" ...
## $ year      : int  1952 1957 1962 1967 1972 1977 1982 1987 1992 1997 ...
## $ lifeExp   : num  28.8 30.3 32 34 36.1 ...
## $ pop       : int  8425333 9240934 10267083 11537966 13079460 14880372 12881816 13867957 16317921 22...
```

So we have 6 columns: country, continent, year, life expectancy, population, GDP per capita.

Look at their types. Some are `chr` meaning they are characters. Others are `int` meaning they are integers, or `numeric` meaning they can also be fractions. For all practical purposes in R, integers and fractions can be used interchangeably, unlike in many languages.

Another option is to use the `head` or `tail` command to look at just the top few or bottom few rows of a dataset.

```
head(gapminder_df)
```

```
##      country continent year lifeExp      pop gdpPercap
## 1 Afghanistan      Asia 1952  28.801  8425333  779.4453
## 2 Afghanistan      Asia 1957  30.332  9240934  820.8530
## 3 Afghanistan      Asia 1962  31.997 10267083  853.1007
## 4 Afghanistan      Asia 1967  34.020 11537966  836.1971
## 5 Afghanistan      Asia 1972  36.088 13079460  739.9811
## 6 Afghanistan      Asia 1977  38.438 14880372  786.1134
```

```
tail(gapminder_df)
```

```
##      country continent year lifeExp      pop gdpPercap
## 1699 Zimbabwe      Africa 1982  60.363  7636524  788.8550
## 1700 Zimbabwe      Africa 1987  62.351  9216418  706.1573
```

```
## 1701 Zimbabwe Africa 1992 60.377 10704340 693.4208
## 1702 Zimbabwe Africa 1997 46.809 11404948 792.4500
## 1703 Zimbabwe Africa 2002 39.989 11926563 672.0386
## 1704 Zimbabwe Africa 2007 43.487 12311143 469.7093
```

Yet another option to “see” the data is the following:

```
View(gapminder_df)
```

The output isn’t shown in the console but is instead shown in a separate tab in RStudio. Helpful, right?

Let’s look at the data for a single year only, viz. 2007.

```
gapminder_2007df <- gapminder_df[gapminder_df$year == 2007,]
```

Let’s look at this much more manageable dataset

```
print(gapminder_2007df)
```

```
##           country continent year lifeExp      pop  gdpPercap
## 12      Afghanistan      Asia  2007  43.828 31889923   974.5803
## 24         Albania      Europe  2007  76.423  3600523  5937.0295
## 36         Algeria      Africa  2007  72.301 33333216  6223.3675
## 48          Angola      Africa  2007  42.731 12420476  4797.2313
## 60        Argentina Americas  2007  75.320 40301927 12779.3796
## 72        Australia Oceania  2007  81.235 20434176 34435.3674
## 84         Austria      Europe  2007  79.829  8199783 36126.4927
## 96         Bahrain      Asia  2007  75.635   708573 29796.0483
## 108      Bangladesh      Asia  2007  64.062 150448339 1391.2538
## 120        Belgium      Europe  2007  79.441 10392226 33692.6051
## 132         Benin      Africa  2007  56.728  8078314 1441.2849
## 144         Bolivia Americas  2007  65.554  9119152  3822.1371
## 156  Bosnia and Herzegovina Europe  2007  74.852  4552198  7446.2988
## 168         Botswana      Africa  2007  50.728  1639131 12569.8518
## 180         Brazil Americas  2007  72.390 190010647  9065.8008
## 192         Bulgaria Europe  2007  73.005  7322858 10680.7928
## 204      Burkina Faso      Africa  2007  52.295 14326203 1217.0330
## 216         Burundi      Africa  2007  49.580  8390505  430.0707
## 228         Cambodia      Asia  2007  59.723 14131858 1713.7787
## 240         Cameroon      Africa  2007  50.430 17696293 2042.0952
## 252         Canada Americas  2007  80.653 33390141 36319.2350
## 264  Central African Republic Africa  2007  44.741  4369038  706.0165
## 276          Chad      Africa  2007  50.651 10238807 1704.0637
## 288         Chile Americas  2007  78.553 16284741 13171.6388
## 300         China      Asia  2007  72.961 1318683096 4959.1149
## 312        Colombia Americas  2007  72.889 44227550 7006.5804
## 324         Comoros      Africa  2007  65.152   710960  986.1479
## 336      Congo, Dem. Rep.      Africa  2007  46.462 64606759  277.5519
## 348         Congo, Rep.      Africa  2007  55.322  3800610 3632.5578
## 360        Costa Rica Americas  2007  78.782  4133884 9645.0614
## 372      Cote d'Ivoire      Africa  2007  48.328 18013409 1544.7501
## 384         Croatia      Europe  2007  75.748  4493312 14619.2227
## 396          Cuba Americas  2007  78.273 11416987 8948.1029
```

## 408	Czech Republic	Europe	2007	76.486	10228744	22833.3085
## 420	Denmark	Europe	2007	78.332	5468120	35278.4187
## 432	Djibouti	Africa	2007	54.791	496374	2082.4816
## 444	Dominican Republic	Americas	2007	72.235	9319622	6025.3748
## 456	Ecuador	Americas	2007	74.994	13755680	6873.2623
## 468	Egypt	Africa	2007	71.338	80264543	5581.1810
## 480	El Salvador	Americas	2007	71.878	6939688	5728.3535
## 492	Equatorial Guinea	Africa	2007	51.579	551201	12154.0897
## 504	Eritrea	Africa	2007	58.040	4906585	641.3695
## 516	Ethiopia	Africa	2007	52.947	76511887	690.8056
## 528	Finland	Europe	2007	79.313	5238460	33207.0844
## 540	France	Europe	2007	80.657	61083916	30470.0167
## 552	Gabon	Africa	2007	56.735	1454867	13206.4845
## 564	Gambia	Africa	2007	59.448	1688359	752.7497
## 576	Germany	Europe	2007	79.406	82400996	32170.3744
## 588	Ghana	Africa	2007	60.022	22873338	1327.6089
## 600	Greece	Europe	2007	79.483	10706290	27538.4119
## 612	Guatemala	Americas	2007	70.259	12572928	5186.0500
## 624	Guinea	Africa	2007	56.007	9947814	942.6542
## 636	Guinea-Bissau	Africa	2007	46.388	1472041	579.2317
## 648	Haiti	Americas	2007	60.916	8502814	1201.6372
## 660	Honduras	Americas	2007	70.198	7483763	3548.3308
## 672	Hong Kong, China	Asia	2007	82.208	6980412	39724.9787
## 684	Hungary	Europe	2007	73.338	9956108	18008.9444
## 696	Iceland	Europe	2007	81.757	301931	36180.7892
## 708	India	Asia	2007	64.698	1110396331	2452.2104
## 720	Indonesia	Asia	2007	70.650	223547000	3540.6516
## 732	Iran	Asia	2007	70.964	69453570	11605.7145
## 744	Iraq	Asia	2007	59.545	27499638	4471.0619
## 756	Ireland	Europe	2007	78.885	4109086	40675.9964
## 768	Israel	Asia	2007	80.745	6426679	25523.2771
## 780	Italy	Europe	2007	80.546	58147733	28569.7197
## 792	Jamaica	Americas	2007	72.567	2780132	7320.8803
## 804	Japan	Asia	2007	82.603	127467972	31656.0681
## 816	Jordan	Asia	2007	72.535	6053193	4519.4612
## 828	Kenya	Africa	2007	54.110	35610177	1463.2493
## 840	Korea, Dem. Rep.	Asia	2007	67.297	23301725	1593.0655
## 852	Korea, Rep.	Asia	2007	78.623	49044790	23348.1397
## 864	Kuwait	Asia	2007	77.588	2505559	47306.9898
## 876	Lebanon	Asia	2007	71.993	3921278	10461.0587
## 888	Lesotho	Africa	2007	42.592	2012649	1569.3314
## 900	Liberia	Africa	2007	45.678	3193942	414.5073
## 912	Libya	Africa	2007	73.952	6036914	12057.4993
## 924	Madagascar	Africa	2007	59.443	19167654	1044.7701
## 936	Malawi	Africa	2007	48.303	13327079	759.3499
## 948	Malaysia	Asia	2007	74.241	24821286	12451.6558
## 960	Mali	Africa	2007	54.467	12031795	1042.5816
## 972	Mauritania	Africa	2007	64.164	3270065	1803.1515
## 984	Mauritius	Africa	2007	72.801	1250882	10956.9911
## 996	Mexico	Americas	2007	76.195	108700891	11977.5750
## 1008	Mongolia	Asia	2007	66.803	2874127	3095.7723
## 1020	Montenegro	Europe	2007	74.543	684736	9253.8961
## 1032	Morocco	Africa	2007	71.164	33757175	3820.1752
## 1044	Mozambique	Africa	2007	42.082	19951656	823.6856

## 1056	Myanmar	Asia	2007	62.069	47761980	944.0000
## 1068	Namibia	Africa	2007	52.906	2055080	4811.0604
## 1080	Nepal	Asia	2007	63.785	28901790	1091.3598
## 1092	Netherlands	Europe	2007	79.762	16570613	36797.9333
## 1104	New Zealand	Oceania	2007	80.204	4115771	25185.0091
## 1116	Nicaragua	Americas	2007	72.899	5675356	2749.3210
## 1128	Niger	Africa	2007	56.867	12894865	619.6769
## 1140	Nigeria	Africa	2007	46.859	135031164	2013.9773
## 1152	Norway	Europe	2007	80.196	4627926	49357.1902
## 1164	Oman	Asia	2007	75.640	3204897	22316.1929
## 1176	Pakistan	Asia	2007	65.483	169270617	2605.9476
## 1188	Panama	Americas	2007	75.537	3242173	9809.1856
## 1200	Paraguay	Americas	2007	71.752	6667147	4172.8385
## 1212	Peru	Americas	2007	71.421	28674757	7408.9056
## 1224	Philippines	Asia	2007	71.688	91077287	3190.4810
## 1236	Poland	Europe	2007	75.563	38518241	15389.9247
## 1248	Portugal	Europe	2007	78.098	10642836	20509.6478
## 1260	Puerto Rico	Americas	2007	78.746	3942491	19328.7090
## 1272	Reunion	Africa	2007	76.442	798094	7670.1226
## 1284	Romania	Europe	2007	72.476	22276056	10808.4756
## 1296	Rwanda	Africa	2007	46.242	8860588	863.0885
## 1308	Sao Tome and Principe	Africa	2007	65.528	199579	1598.4351
## 1320	Saudi Arabia	Asia	2007	72.777	27601038	21654.8319
## 1332	Senegal	Africa	2007	63.062	12267493	1712.4721
## 1344	Serbia	Europe	2007	74.002	10150265	9786.5347
## 1356	Sierra Leone	Africa	2007	42.568	6144562	862.5408
## 1368	Singapore	Asia	2007	79.972	4553009	47143.1796
## 1380	Slovak Republic	Europe	2007	74.663	5447502	18678.3144
## 1392	Slovenia	Europe	2007	77.926	2009245	25768.2576
## 1404	Somalia	Africa	2007	48.159	9118773	926.1411
## 1416	South Africa	Africa	2007	49.339	43997828	9269.6578
## 1428	Spain	Europe	2007	80.941	40448191	28821.0637
## 1440	Sri Lanka	Asia	2007	72.396	20378239	3970.0954
## 1452	Sudan	Africa	2007	58.556	42292929	2602.3950
## 1464	Swaziland	Africa	2007	39.613	1133066	4513.4806
## 1476	Sweden	Europe	2007	80.884	9031088	33859.7484
## 1488	Switzerland	Europe	2007	81.701	7554661	37506.4191
## 1500	Syria	Asia	2007	74.143	19314747	4184.5481
## 1512	Taiwan	Asia	2007	78.400	23174294	28718.2768
## 1524	Tanzania	Africa	2007	52.517	38139640	1107.4822
## 1536	Thailand	Asia	2007	70.616	65068149	7458.3963
## 1548	Togo	Africa	2007	58.420	5701579	882.9699
## 1560	Trinidad and Tobago	Americas	2007	69.819	1056608	18008.5092
## 1572	Tunisia	Africa	2007	73.923	10276158	7092.9230
## 1584	Turkey	Europe	2007	71.777	71158647	8458.2764
## 1596	Uganda	Africa	2007	51.542	29170398	1056.3801
## 1608	United Kingdom	Europe	2007	79.425	60776238	33203.2613
## 1620	United States	Americas	2007	78.242	301139947	42951.6531
## 1632	Uruguay	Americas	2007	76.384	3447496	10611.4630
## 1644	Venezuela	Americas	2007	73.747	26084662	11415.8057
## 1656	Vietnam	Asia	2007	74.249	85262356	2441.5764
## 1668	West Bank and Gaza	Asia	2007	73.422	4018332	3025.3498
## 1680	Yemen, Rep.	Asia	2007	62.698	22211743	2280.7699
## 1692	Zambia	Africa	2007	42.384	11746035	1271.2116

```
## 1704 Zimbabwe Africa 2007 43.487 12311143 469.7093
```

Now let's find the average life expectancy (in 2007) of countries in the various continents

Step 1. Filter by each continent. Let's start with Asia.

```
print(gapminder_2007df[gapminder_2007df$continent == "Asia",])
```

```
##      country continent year lifeExp      pop  gdpPercap
## 12  Afghanistan      Asia  2007  43.828  31889923    974.5803
## 96   Bahrain        Asia  2007  75.635   708573  29796.0483
## 108  Bangladesh      Asia  2007  64.062  150448339   1391.2538
## 228   Cambodia       Asia  2007  59.723   14131858   1713.7787
## 300    China         Asia  2007  72.961  1318683096   4959.1149
## 672 Hong Kong, China  Asia  2007  82.208   6980412  39724.9787
## 708    India         Asia  2007  64.698  1110396331   2452.2104
## 720   Indonesia     Asia  2007  70.650  223547000    3540.6516
## 732    Iran         Asia  2007  70.964   69453570  11605.7145
## 744    Iraq         Asia  2007  59.545   27499638   4471.0619
## 768    Israel       Asia  2007  80.745   6426679   25523.2771
## 804    Japan        Asia  2007  82.603  127467972   31656.0681
## 816    Jordan       Asia  2007  72.535   6053193   4519.4612
## 840 Korea, Dem. Rep.  Asia  2007  67.297   23301725   1593.0655
## 852    Korea, Rep.   Asia  2007  78.623   49044790   23348.1397
## 864    Kuwait       Asia  2007  77.588   2505559   47306.9898
## 876    Lebanon      Asia  2007  71.993   3921278   10461.0587
## 948    Malaysia     Asia  2007  74.241   24821286  12451.6558
## 1008   Mongolia     Asia  2007  66.803   2874127   3095.7723
## 1056   Myanmar      Asia  2007  62.069   47761980    944.0000
## 1080   Nepal        Asia  2007  63.785   28901790   1091.3598
## 1164    Oman        Asia  2007  75.640   3204897   22316.1929
## 1176   Pakistan     Asia  2007  65.483  169270617   2605.9476
## 1224   Philippines   Asia  2007  71.688   91077287   3190.4810
## 1320   Saudi Arabia   Asia  2007  72.777   27601038  21654.8319
## 1368    Singapore    Asia  2007  79.972   4553009   47143.1796
## 1440    Sri Lanka     Asia  2007  72.396   20378239   3970.0954
## 1500    Syria        Asia  2007  74.143   19314747   4184.5481
## 1512    Taiwan       Asia  2007  78.400   23174294  28718.2768
## 1536    Thailand     Asia  2007  70.616   65068149   7458.3963
## 1656    Vietnam     Asia  2007  74.249   85262356   2441.5764
## 1668 West Bank and Gaza Asia  2007  73.422   4018332   3025.3498
## 1680    Yemen, Rep.   Asia  2007  62.698   22211743   2280.7699
```

Here, as you see, only those countries that are in Asia are shown.

Step 2. Fetch the life expectancy of Asian countries in 2007.

```
print(gapminder_2007df[gapminder_2007df$continent == "Asia",]$lifeExp)
```

```
## [1] 43.828 75.635 64.062 59.723 72.961 82.208 64.698 70.650 70.964 59.545
## [11] 80.745 82.603 72.535 67.297 78.623 77.588 71.993 74.241 66.803 62.069
## [21] 63.785 75.640 65.483 71.688 72.777 79.972 72.396 74.143 78.400 70.616
## [31] 74.249 73.422 62.698
```


Step 3. Calculate the mean.

```
print(mean(gapminder_2007df[gapminder_2007df$continent == "Asia",]$lifeExp))
```

```
## [1] 70.72848
```

Repeat for the other continents. Which continent had the highest average life expectancy in 2007?


Hint: you can find the names of all the continents using the `unique` command like this:

```
print(unique(gapminder_2007df$continent))
```

```
## [1] "Asia"      "Europe"    "Africa"    "Americas"  "Oceania"
```

Finally, let's draw our first R graph. Let's have "life expectancy on X axis" and "GDP per capita" on the y axis.

```
plot(gapminder_2007df$lifeExp, gapminder_2007df$gdpPercap)
```



Now, let's draw a "regression line" between these two variables: life expectancy and GDP per capita. To do that, first plot the graph, and then draw the line using the `abline` command. Notice the `lm` command and the use of the tilde `~`. This is instructing R to fit a "linear model" (`lm`) between GDP per capita and life expectancy.

```
plot(gapminder_2007df$lifeExp, gapminder_2007df$gdpPercap)
abline(lm(gapminder_2007df$gdpPercap ~ gapminder_2007df$lifeExp))
```



Exercises:

Using what you have learned so far, answer the following questions using the gapminder dataset.

1. Which country had the 5th highest population in 1957? (Hint: use the `order` command)
2. Which continent had the lowest average population in 1982?
3. Which country had the highest per capita GDP in 1952?